

# SDR KT Onboarding Agent — Process Redesign

MGT 449 — Milestone 04 — Group 2

Anjali Roy


Shreya Anekar

2026-06-04

## AI-supported SDR onboarding workflow

The diagram below maps the redesigned process. Color legend follows the Milestone 01 color system:

- **Green** — Human task
- **Red** — AI task (GenAI)
- **Purple** — Human review or approval
- **Orange** — Human + AI together
- **Blue** — Regular automation (no GenAI)



## What changed from the original process

**Milestone 01 baseline.** A new SDR was onboarded manually by enablement (~2 hours per hire of customization), ramped through a generic Trailhead checklist regardless of prior experience, coached ad-hoc by their manager (3–5 hours per week of unstructured call review, 24–72 hour feedback latency), and certified through an opaque process by week 12. Once certified, there was no structured way to capture what worked (or didn’t) during the ramp, so the onboarding program rarely improved cohort over cohort. Per RepVue Feb 2026, 42% of new SDRs were hitting quota.

**Redesigned workflow.** Three structural changes:

1. **An orchestrated pre-boarding chain** drafts the plan, ranks a mentor, and writes the welcome email before Day 1. A single review gate gives the manager an approve/edit/reject/escalate decision with the cited sources, the confidence score, and an explicit warning when confidence is low.
2. **Five on-demand AI coworker features** carry the SDR through the 12-week ramp and out the door:
  - **AI Tutor** — ChromaDB-semantic retrieval over 16 KB docs with inline [n] citations and SDR-side approve/reject.
  - **Cold Call Simulation** — stateless persona LLM + scored debrief on the same Rubric v3.1 used by the coaching scorer.
  - **Coaching Notes** — LLM scores a pasted call transcript against the rubric and drafts a structured note; manager reviews before delivery.
  - **Certification Gap Analysis** — LLM scores the SDR’s simulation transcripts and certification exam against a 5-dimension rubric with quoted evidence per dimension; a server-side band recompute overrides the LLM if its self-reported verdict disagrees with the dimension scores; the trainer makes the final Pass/Fail with a required comment.
  - **Onboarding Debrief** — after passing certification the SDR picks their own name from a dropdown (auto-fills territory/vertical/experience) and walks a 7-question debrief on their 12-week ramp; the AI synthesizes a manager-facing debrief (per-element feedback, ramp barriers, suggested program improvements); the manager reviews and shares it with enablement.
3. **Auditability everywhere.** Every approve / edit / reject / escalate decision is appended to `data/decision_log.json` with the feature, action, and context. The dashboard reads this log.

## Which steps should be faster or easier

Per `estimates.md` (all assumptions, ranges labeled):

| Step                              | Current                   | AI-supported     |
|-----------------------------------|---------------------------|------------------|
| Build pre-boarding plan           | 90–150 min                | 10–20 min        |
| Rank Trail Guide mentor           | 20–40 min                 | 5–10 min         |
| Draft welcome email               | 15–30 min                 | 3–8 min          |
| Answer a product/policy question  | 10–30 min                 | 1–3 min          |
| Coaching note from a call         | 20–40 min                 | 5–12 min         |
| <b>Certification gap analysis</b> | <b>45–75 min</b>          | <b>10–15 min</b> |
| Post-cert onboarding debrief      | 60–90 min (often skipped) | 10–15 min review |

Net effect on the pre-boarding cycle: 125–220 min → 18–38 min; roughly \$120–215 saved per hire at \$70/hr fully-loaded.

## Which steps still require human judgment

Five explicit review gates remain:

1. **Manager gate** at the end of the pre-boarding chain — approve, edit, reject, escalate.
2. **SDR-side approve/reject** on every tutor answer (the rejection note is logged so future retrieval can avoid the bad chunk).
3. **Manager review** of every coaching note before it reaches the SDR.
4. **Trainer decision** on certification — required comment, AI recommendation is advisory only, server-recomputed band is the source of truth.
5. **Manager review** of every onboarding debrief before it is shared with the enablement team.

The three Milestone 01 checkpoints that are legally, financially, or organizationally loaded — certification, AE/territory assignment, and qualified-meeting verification — stay fully human.

## Which steps are risky if fully automated

- **Certification verdict drift.** The LLM might self-report PASS while dimension scores warrant BORDERLINE. **Mitigation:** server-side band recompute overrides; mismatch added to `review_reasons` so the trainer sees *why* it's flagged.
- **Onboarding plan hallucination.** The LLM might invent a Trailhead module. **Mitigation:** module-ID allow-list enforced against the catalog before the artifact is written.
- **Coaching note unfairness.** A bad transcript or miscalibrated rubric could deliver demoralizing feedback. **Mitigation:** required manager approval before delivery.
- **Tutor fabrication.** The May 25 baseline DeepEval runs showed the tutor adding unsupported claims in 2 of 10 cases. **Mitigation:** the tightened prompt's "block ungrounded identifiers" rule + Faithfulness check in the eval suite caught one of those (case 10 module-ID fabrication moved from 1/3 to 2/3 passing post-fix).
- **Debrief over-generalization.** AI might over- or understate a program issue from a sparse answer. **Mitigation:** low confidence flags review and the manager edits the affected section before approving.

## What the AI coworker needs from data, tools, or documents

**Data sources** (under `data/`):

- `raw/sdr_profiles.csv` — new-hire records
- `raw/ae_profiles.csv` — candidate mentors
- `raw/summaries/` — per-hire summaries
- `raw/sim_personas.json` — prospect personas for the cold-call simulator
- `sdr_records/` — 7-SDR synthetic cohort (simulation transcripts, exam answers, coaching notes, cached gap analyses)
- Runtime transcripts pasted into Coaching Notes; debrief answers from the Onboarding Debrief

**Documents** — 16 docs in `rag/knowledge_base/`, including the certification rubric and the canonical 5-question exam bank for the cert feature, plus 7 scraped Wikipedia refs as background context.

**Tools / Python functions** — 7 LLM skills, 2 deterministic automations, 2 MCP servers, 2 retrievers (BM25 + ChromaDB), plus the FastAPI orchestrator (used by the UI) and the CLI orchestrator (used for offline chain runs).

## **Final business output**

**A trained, certified, ramped SDR plus a structured debrief that feeds back into the onboarding program** — produced in 18–38 minutes of manager time per pre-boarding cycle instead of 125–220 minutes, with explicit human review at five gates and an audit log of every decision.