

SDR KT Onboarding Agent

An AI coworker for the 12-week SDR journey — with a human at every gate

MGT 449 · Milestone 04 · Group 2 — Anjali Roy & Shreya Anekar

2026-06-04

The big idea

ONBOARDING AN SDR TODAY

A new Sales Development Rep takes **12 weeks** to ramp — and today that ramp is slow, inconsistent, and barely measured.

42%

hit quota

new SDRs (RepVue, Feb 2026)

24–72h

coaching feedback lag

on cold-call reviews

3–5h


manual call review

per manager, per week

Our redesign — one AI coworker for the whole journey: draft the plan → coach the calls → certify at week 12 → debrief to improve the program. A **Manager** reviews every consequential step.

The redesigned workflow

Four phases — pre-boarding, ramp, certification, post-cert debrief — with a Manager review gate in each.



Live demo — two SDRs, two paths WATCH IT RUN

Priya — the clean path

- Entry SDR · EMEA · healthcare
- AI drafts a personalized 4-week plan + cited sources
- High confidence → **Manager approves** at the gate

Marcus — the review path

- Mid SDR · AMER-West · SaaS
- AI flags **low confidence** (missing profile data)
- Run pauses → **Manager reviews** before anything ships

Real input → useful output, with the **evidence visible** and a **Manager decision** at each step.

Under the hood

Every request — a hire, a question, a transcript — flows through the same pipeline.



One uniform envelope between every step makes the workflow **auditable, idempotent, and safe to gate**.

Data, evidence & AI usage

WHAT THE AI USES — AND WHAT THE USER CAN INSPECT

Data it uses

SDR & AE profiles · per-hire summaries · simulation personas · a 7-SDR records set · call transcripts pasted by the user.

Documents (RAG)

A **16-doc** internal knowledge base — onboarding policy, quota ramp, objection handbook, cold-call rubric, territory + vertical playbooks, certification rubric & exam — plus 7 public references.

Tools / functions

7 LLM skills · 2 deterministic automations · 2 retrievers (BM25 + ChromaDB) · 2 MCP servers.

Evidence the user sees

An **EvidencePanel** with numbered cited sources + match scores · inline [n] citations that open the exact passage · per-dimension cards quoting the SDR's own work verbatim.

When evidence is weak or missing: the tutor cites a source or refuses (no bluffing) · confidence < 0.60 blocks auto-complete and pauses for review · a server-side recompute overrides the AI's certification verdict if they disagree.

Human stays in control — five gates

- 1 **Pre-boarding plan** — Manager approves / edits / rejects / escalates
- 2 **AI Tutor answer** — SDR approves / rejects; rejection re-enters retrieval
- 3 **Coaching note** — Manager reviews before the SDR sees it
- 4 **Certification** — Manager decides Pass/Fail with a required comment
- 5 **Onboarding debrief** — Manager reviews, then shares with enablement

Two hard safeguards

Confidence < 0.60 **blocks auto-complete** and pauses for review · a server-side recompute **overrides** the AI's certification verdict if they disagree.

Every decision logged

Approve / edit / reject / escalate is written to `decision_log.json` with feature, action and timestamp — a full audit trail.

The business case

TIME, COST & QUALITY · \$70/HR FULLY-LOADED (ASSUMPTION)

Pre-boarding step	Today	With AI
Build plan	90–150 min	10–20 min
Match mentor	20–40 min	5–10 min
Welcome email	15–30 min	3–8 min
Total per hire	125–220 min	18–38 min
Cost per hire	\$146–257	\$21–44

\$120–215
saved per hire
pre-boarding alone

Coaching review
3–5 hr/wk → **under 1 hr/wk** of manual call review.

Certification
45–75 → **10–15 min** per SDR (\$40–71 saved each).

Quality gains
Consistent rubric coaching · cited answers · faster feedback · every decision logged.

Testing — the kinds of checks we ran

We didn't just ship the AI Tutor — an automated **DeepEval** suite judges its answers, rerun to trust the trend over noise.

10 cases × 4 checks
Faithfulness — stays within retrieved sources? · **Answer relevancy** — answers what was asked? · **Citation + next-step** — cites [n] and a concrete action? · **SDR face-validity** — useful to a real SDR? Cases span everyday, edge (vague / no-source), and predicted-failure inputs.

What it caught — and the fix
It flagged real failures: over-answering vague questions, a fabricated module-ID, generic next-steps. We changed **one thing** (narrow broad questions + block ungrounded identifiers), reran 6× — the weak case went from failing to passing and fabrication disappeared.

Honest: one rule over-formalized a couple of answers, and the biggest remaining gap is **knowledge-base coverage** — a content fix, not a prompt fix.

Face validity — would a sales leader believe this?

The size is believable

The AI only compresses what it's genuinely good at — drafting plans, retrieving answers, scoring calls. The Manager's review time is **added back in**, so the headline cuts are net savings, not magic.

The direction is right

Every consequential decision still needs a human: the Manager approves plans and coaching notes, and the certification AI is **advisory** — the Manager makes the Pass/Fail call.

It's grounded, not guessed

Tutor answers cite the source passage behind each claim (Faithfulness ~8.7/10 across our runs). When the AI isn't sure, it **flags for review** instead of bluffing.

Where we're honest

Certification is demoed on a synthetic 7-SDR cohort, and the recurring coaching saving assumes a typical weekly call volume — both clearly labeled as assumptions.

What we'd improve

Make the AI sharper

- Align knowledge-base docs to the test rubrics (the main remaining failure)
- Soften the rule that over-formalized two answers
- Capture real certification artifacts from the live cold-call sim

Make it production-ready

- Persist state to a real database (today: in-memory)
- Auth + role enforcement at the API layer
- CI running tests + evals on every change

Thank you

SDR KT Onboarding Agent — an AI coworker with a human at every gate.

Anjali Roy · Shreya Anekar · Group 2

Repo · `rsm-genai-2026/mgt449-milestone04-Group-2` Demo · `uv run python workflow_server.py` + `cd ui && npm run dev`

Questions?